

# ADS 503 Final Project

Claire Phibbs, Christopher Richardson, Martin Zagari

2022-06-17

## The Data

Our target variable “Risk Performance” along with the first 13 predictor columns

```
file_loc <- '/Volumes/GoogleDrive/My Drive/503/Project 503/Fico Data/heloc_dataset_v1.csv'
heloc <- read.csv(file_loc)

# sort col names for readability purposes
heloc <- heloc[, order(names(heloc))]


knitr::kable(heloc[1:4,c(24,1:13)]) %>%
  kableExtra::kable_styling("striped", full_width = F) %>%
  kableExtra::row_spec(0, angle = -90)
```

Predictor columns 14 to 23

```
knitr::kable(heloc[1:4,c(24,14:23)]) %>%
  kableExtra::kable_styling("striped", full_width = F) %>%
  kableExtra::row_spec(0, angle = -90)
```

## Data Pre-Processing


```
DataExplorer::plot_intro(heloc)
```



[illegible]


PercentTradesWBalance		69
PercentTradesNeverDelq	83	0
PercentInstallTrades	43	67
NumTradesOpeninLast12M	1	4
NumTrades90Ever2DerogPubRec	0	0
NumTrades60Ever2DerogPubRec	3	4
NumTotalTrades	23	7
NumSatisfactoryTrades	20	9
NumRevolvingTradesWBalance	8	4
NumInstallTradesWBalance	1	2
RiskPerformance	Bad	Bad

```
# -9 = No Credit History
# -8 and -7 = No recent activity
heloc[heloc == -9] <- NA
heloc[heloc == -8] <- NA
heloc[heloc == -7] <- NA
DataExplorer::plot_intro(heloc)
```




appears that roughly 76.1% of our data had one or more predictors had a -9,-8,-7.

```
# removing missing values from rows that span across all columns (588 values)
heloc_No_NA <- heloc %>% dplyr::filter_at(vars(-RiskPerformance),
                                          any_vars(!is.na(.)))
DataExplorer::plot_intro(heloc_No_NA)
```



```
# bar plot of response variable; RiskPerformance
barplot(table(heloc_No_NA$RiskPerformance),
        main="Plot of Response Variable: RiskPerformance",
        xlab="RiskPerformance")
```

## Plot of Response Variable: RiskPerformance



```
table(heloc_No_NA$RiskPerformance)
```

```
##
##  Bad Good
## 5136 4735
```

A 51:47 split! Nearly a 50:50 balance!

```
# create training indices
set.seed(3)
heloc_training <- caret::createDataPartition(heloc_No_NA$RiskPerformance,
                                              p=0.8,
                                              list=FALSE)

# training/set sets
heloc_train <- heloc_No_NA[heloc_training, ]
heloc_test  <- heloc_No_NA[-heloc_training, ]

# knn imputation
heloc_impute <- caret::preProcess(heloc_train,
                                  method = 'knnImpute')

heloc_train <- stats::predict(heloc_impute,
                              newdata=heloc_train)

heloc_test  <- stats::predict(heloc_impute,
                              newdata=heloc_test)

# remove highly correlated predictors
high_corr <- caret::findCorrelation(stats::cor(heloc_train[, -24]),
                                    0.85)

# removal of high cor predictors
heloc_train <- heloc_train[, -(high_corr)]
heloc_test  <- heloc_test[, -(high_corr)]
names(heloc_test[high_corr])
```

```
## [1] "NumTradesOpeninLast12M" "NumInstallTradesWBalance"
```

We keep the heloc\_train/test dataframes for formula based functions.

```

no_risk_bool <- names(heloc_train) != 'RiskPerformance'

# x = predictors
heloc_train_x <- heloc_train[,no_risk_bool]
heloc_test_x <- heloc_test[,no_risk_bool]

# y = response/target
heloc_train_y <- heloc_train[, 'RiskPerformance']
heloc_test_y <- heloc_test[, 'RiskPerformance']


```

## EDA

```

# correlations
png(file="Corr_Plot.png")
# correlations
corrplot::corrplot(stats::cor(heloc_train),
                    number.cex = 0.5,
                    tl.cex = 0.8)
dev.off()

```




Correlation Plot

Outcome Bar Plot

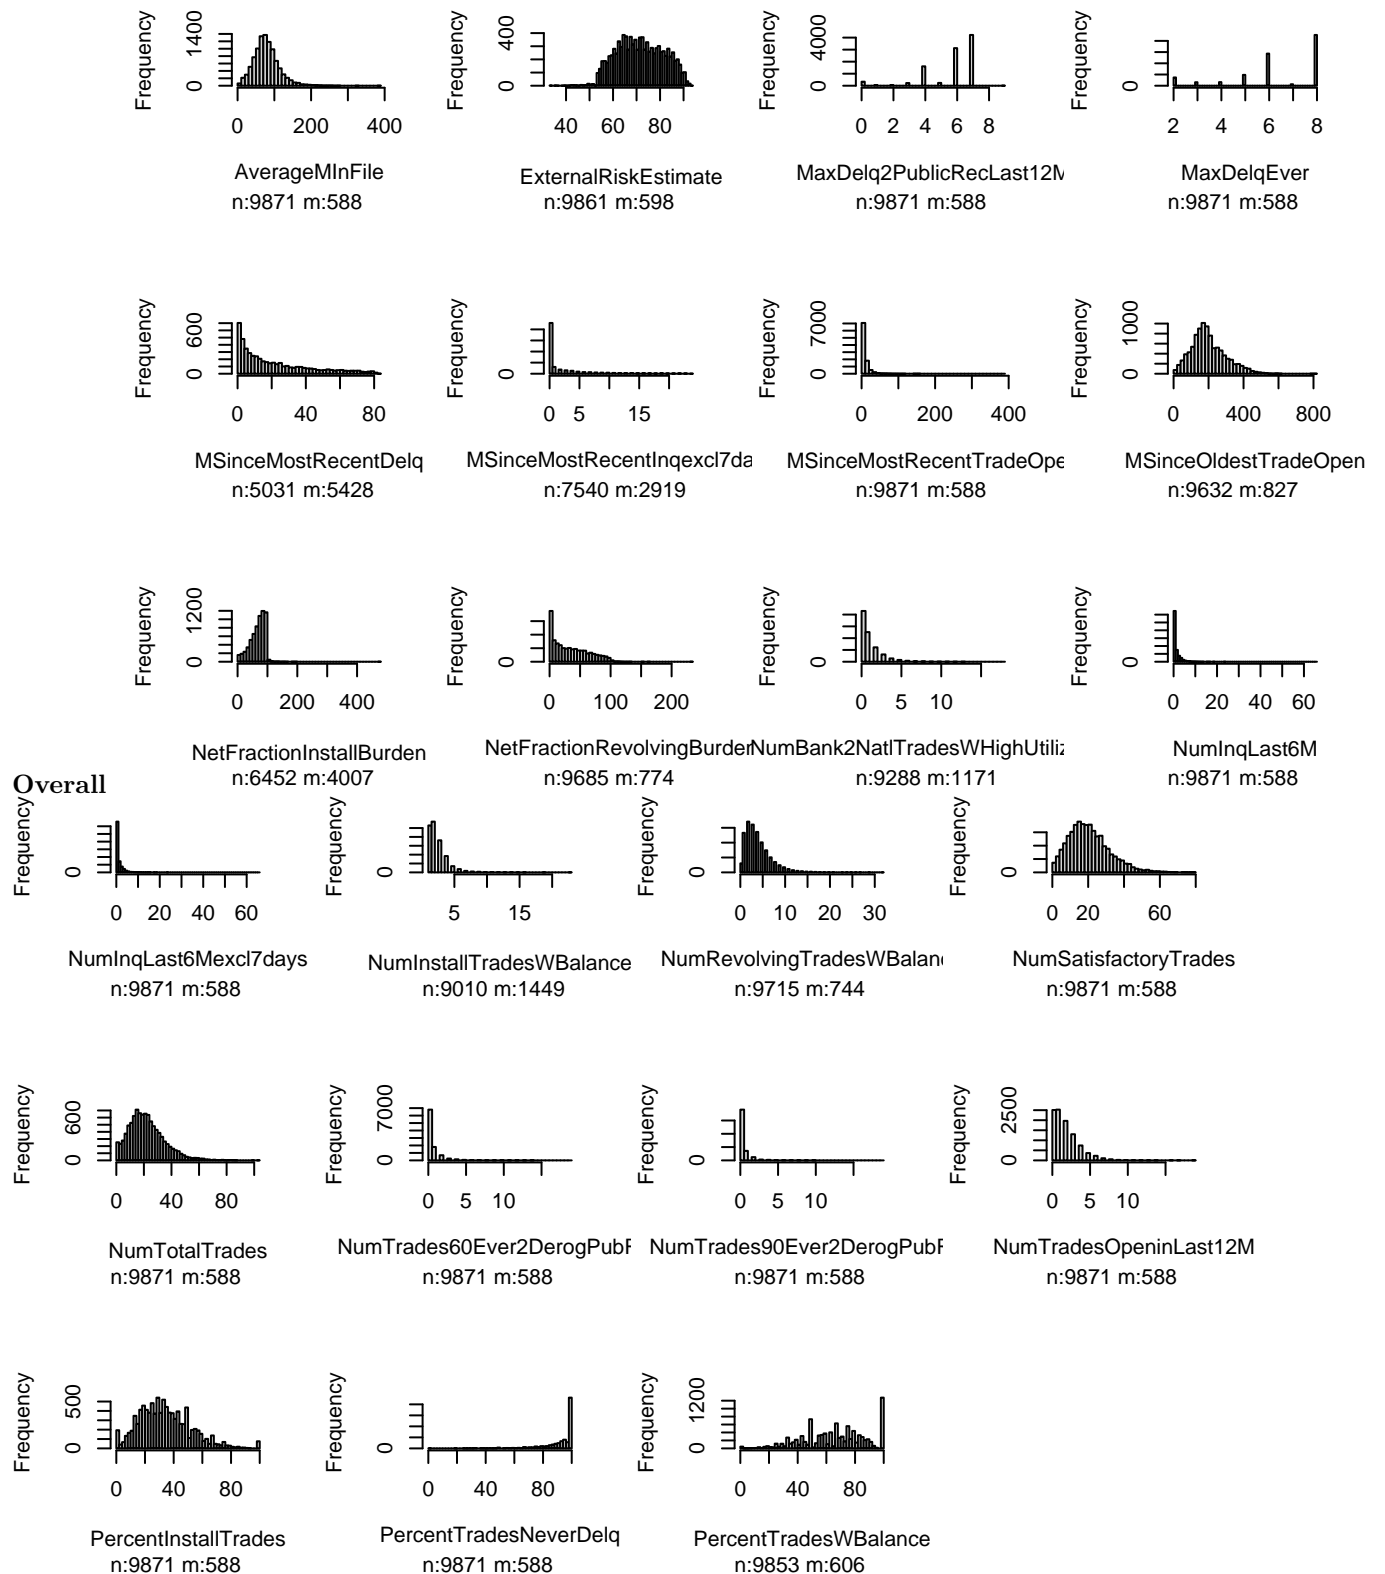
```
# bar plot of response variable; RiskPerformance
barplot(table(heloc$RiskPerformance), main="Plot of Response Variable: RiskPerformance", xlab="RiskPerformance")
```

## Plot of Response Variable: RiskPerformance



### Histograms

```
# histograms to view predictor variable frequencies  
par(mfrow=c(3,4))  
Hmisc::hist.data.frame(heloc[,1:23])
```



#### By Outcome

```
# Deselect Bool Outcome/Response/Target variable
no_risk_bool <- names(heloc_train) != 'RiskPerformance'
```




```

# heloc_imputed_full_set <- dplyr::as_tibble(heloc_imputed_full_set)
heloc_ifs_names <- colnames(heloc_train[,no_risk_bool])


# empty list to gather all
plot_list <- list()
for (name in heloc_ifs_names){
  p <- heloc_train %>%
    ggplot( aes(x=heloc_train[,name], fill=RiskPerformance)) +
    geom_histogram( color="#e9ecef", alpha=0.6, position = 'identity') +
    scale_fill_manual(values=c("#69b3a2", "#404080")) +
    xlab(name) +
    labs(fill="")
  print(p)
}

```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




## `stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.



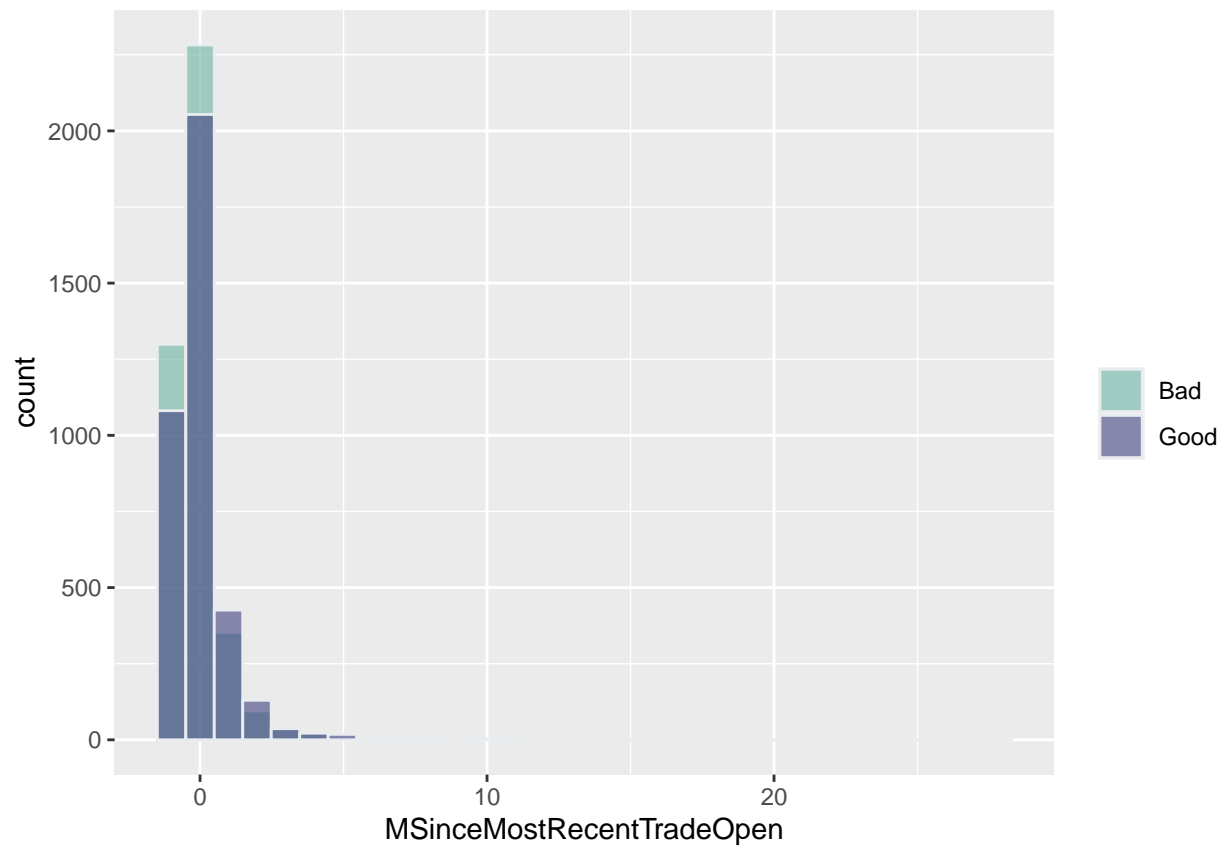
```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




## `stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.



```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```







```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




## `stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



## `stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```







```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




## BoxPlots


```
# boxplot to view outliers
boxplot(heloc[, 1:23])
```





```
for (name in heloc_ifs_names){  
  # boxplot to view outliers  
  p <- heloc_train %>%  
    ggplot( aes(x=heloc_train[,name], y=name, fill=RiskPerformance)) +  
    geom_boxplot() +  
    xlab(name)  
  print(p)  
}
```







By Group








































```
control <- caret::trainControl(method="cv",
                                classProbs=TRUE,
                                savePredictions=TRUE,
                                summaryFunction=twoClassSummary)
```

## Models

With the following models, we are trying to maximize the SPECIFICITY of the models to prevent any non-qualified loanees being presented with a HELOC loan.

Specificity is defined as: “The specificity is defined as the rate that nonevent samples are predicted as nonevents” (Kuhn & Johnson, 2013)

```
bestIndex <- function(model){
  # returns top ROC value and surrounding indices from model
  highest_score <- max(model$results$ROC)

  # get row index and convert to type Int
  best_index <- rownames(model$results[model$results$ROC == highest_score,])
  best_index <- as.integer(best_index)

  return(best_index)
}

confusionMatrix <- function(testResults.model){
  caret::confusionMatrix(testResults.model,
                          as.factor(testResults$obs),
```

```

        positive="Good")
}

modelScoreBoard <- function(testResults){
  # feed in testResults dataframe and out comes a model scoreboard!
  scoreboard <- data.frame()

  bool <- names(testResults) != 'obs'
  col_names <- colnames(testResults[,bool])

  for (colname in col_names){

    testResults.model <- testResults[,colname]
    cf <- caret::confusionMatrix(testResults.model,
                                as.factor(testResults$obs),
                                positive="Good")

    # gather testResults
    acc <- data.frame(metric=cf$overall[1])
    # gather Precision, Sensitivity, Specificity, & F1
    metrics <- list(cf$byClass[c(5,1,2,7)])
    metrics <- data.frame(Metrics=metrics)
    names(metrics) <- 'metric'
    # gather all metrics in 1 df
    metrics <- rbind(acc,metrics)
    names(metrics) <- colname

    metrics <- t(metrics)
    scoreboard <- rbind(scoreboard, metrics)
  }
  return(scoreboard)
}

# helper function for roc
roc_build <- function(model) {
  THE_ROC <- roc(response = model$pred$obs,
                 predictor = model$pred$Bad,
                 levels = rev(levels(model$pred$obs)))
  return(THE_ROC)
}

```

## Discriminant Classification Models

### LDA

```

set.seed(100)
lda_model <- caret::train(x=heloc_train_x,
                          y=heloc_train_y,
                          method="lda",
                          metric="ROC",
                          trControl=control)

lda_modelRoc <- roc_build(lda_model)

```

```

## Setting direction: controls < cases
lda_model

## Linear Discriminant Analysis
##
## 7897 samples
## 21 predictor
## 2 classes: 'Bad', 'Good'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 7107, 7107, 7107, 7108, 7107, 7108, ...
## Resampling results:
##
## ROC      Sens      Spec
## 0.7998689 0.7527435 0.7114538

lda_predictions <- stats::predict(lda_model, heloc_test_x)

# create dataframe to store
testResults <- data.frame(obs=heloc_test_y,
                           lda_model=lda_predictions)

# confusion matrix
confusionMatrix(testResults$lda_model)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction Bad Good
##      Bad 768 288
##      Good 259 659
##
##              Accuracy : 0.7229
##              95% CI : (0.7026, 0.7426)
##      No Information Rate : 0.5203
##      P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 0.4442
##
## Mcnemar's Test P-Value : 0.2312
##
##              Sensitivity : 0.6959
##              Specificity : 0.7478
##      Pos Pred Value : 0.7179
##      Neg Pred Value : 0.7273
##      Prevalence : 0.4797
##      Detection Rate : 0.3338
##      Detection Prevalence : 0.4650
##      Balanced Accuracy : 0.7218
##
##      'Positive' Class : Good
##

```

## Logistic Regression

```
set.seed(100)
logreg_model <- caret::train(x=heloc_train_x,
                             y=heloc_train_y,
                             method="glm",
                             metric="ROC",
                             trControl=control)

testResults$log_reg_model <- stats::predict(logreg_model, heloc_test_x)
logreg_modelRoc <- roc_build(logreg_model)
```

```
## Setting direction: controls < cases
```

```
logreg_model
```

```
## Generalized Linear Model
```

```
##
```

```
## 7897 samples
```

```
## 21 predictor
```

```
## 2 classes: 'Bad', 'Good'
```

```
##
```

```
## No pre-processing
```

```
## Resampling: Cross-Validated (10 fold)
```

```
## Summary of sample sizes: 7107, 7107, 7107, 7108, 7107, 7108, ...
```

```
## Resampling results:
```

```
##
```

```
## ROC Sens Spec
```

```
## 0.8000842 0.7554216 0.7130383
```


```
logreg_model$finalModel$coefficients
```

```
## (Intercept) AverageMinFile
## -0.1803777116 0.2470129097
## ExternalRiskEstimate MaxDelq2PublicRecLast12M
## 0.4763324385 0.0454362796
## MaxDelqEver MSinceMostRecentDelq
## 0.0320452303 0.1630765084
## MSinceMostRecentInqexcl7days MSinceMostRecentTradeOpen
## 0.2726958674 -0.0650335190
## MSinceOldestTradeOpen NetFractionInstallBurden
## 0.0649987498 -0.0543543815
## NetFractionRevolvingBurden NumBank2NatlTradesWHighUtilization
## -0.2929249411 -0.1376246387
## NumInqLast6M NumInstallTradesWBalance
## -0.1638154810 0.0167061783
## NumRevolvingTradesWBalance NumSatisfactoryTrades
## -0.2278121804 0.3689635005
## NumTotalTrades NumTrades90Ever2DerogPubRec
## 0.0194119014 0.0001783062
## NumTradesOpeninLast12M PercentInstallTrades
## -0.0196622137 -0.1878716750
## PercentTradesNeverDelq PercentTradesWBalance
## 0.1716749452 0.1173618900
```

```
confusionMatrix(testResults$log_reg_model)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction Bad Good
##      Bad 766 288
##      Good 261 659
##
##           Accuracy : 0.7219
##           95% CI : (0.7015, 0.7416)
##      No Information Rate : 0.5203
##      P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.4422
##
##  Mcnemar's Test P-Value : 0.2671
##
##           Sensitivity : 0.6959
##           Specificity : 0.7459
##      Pos Pred Value : 0.7163
##      Neg Pred Value : 0.7268
##           Prevalence : 0.4797
##      Detection Rate : 0.3338
##      Detection Prevalence : 0.4661
##      Balanced Accuracy : 0.7209
##
##      'Positive' Class : Good
##
```

```
lr_varImp <- caret::varImp(logreg_model, scale=FALSE)
plot(lr_varImp, top=20)
```



####

#### Cost Matrix Threshold Analysis

```
#get raw probs from model
predictions <- predict(logreg_model, heloc_test_x, type = 'prob')
predictions$OBS <- as.factor(heloc_test$RiskPerformance)
predictions <- predictions %>%
  mutate(lr10 = as.factor(if_else(Bad > 0.1, 'Bad', 'Good')))
predictions <- predictions %>%
  mutate(lr20 = as.factor(if_else(Bad > 0.2, 'Bad', 'Good')))
predictions <- predictions %>%
  mutate(lr30 = as.factor(if_else(Bad > 0.3, 'Bad', 'Good')))
predictions <- predictions %>%
  mutate(lr40 = as.factor(if_else(Bad > 0.4, 'Bad', 'Good')))
predictions <- predictions %>%
  mutate(lr50 = as.factor(if_else(Bad > 0.5, 'Bad', 'Good')))
predictions <- predictions %>%
  mutate(lr60 = as.factor(if_else(Bad > 0.6, 'Bad', 'Good')))
predictions <- predictions %>%
  mutate(lr70 = as.factor(if_else(Bad > 0.7, 'Bad', 'Good')))
predictions <- predictions %>%
  mutate(lr80 = as.factor(if_else(Bad > 0.8, 'Bad', 'Good')))
predictions <- predictions %>%
  mutate(lr90 = as.factor(if_else(Bad > 0.9, 'Bad', 'Good')))

# cf function
cost_confusionMatrix <- function(prediction.rate){
  cm <- caret::confusionMatrix(prediction.rate,
                                predictions$OBS,
                                positive = "Bad")

  return(cm)
}
```

```

CF10 <- cost_confusionMatrix(predictions$lr10)
CF20 <- cost_confusionMatrix(predictions$lr20)
CF30 <- cost_confusionMatrix(predictions$lr30)
CF40 <- cost_confusionMatrix(predictions$lr40)
CF50 <- cost_confusionMatrix(predictions$lr50)
CF60 <- cost_confusionMatrix(predictions$lr60)
CF70 <- cost_confusionMatrix(predictions$lr70)
CF80 <- cost_confusionMatrix(predictions$lr80)
CF90 <- cost_confusionMatrix(predictions$lr90)

```

```

Costs = matrix(c(0,-1000*.85,-60,60), ncol=2, nrow=2)

```

```

Prev = matrix(c(9.6/50,9.6/50,2,2), ncol=2, nrow=2)

```

```

CF10$table

```

```

##           Reference
## Prediction Bad Good
##           Bad 1023 895
##           Good   4  52

```

```

sum(CF10$table*Costs*Prev)

```

```

## [1] -101812.8

```

```

CF20$table

```

```

##           Reference
## Prediction Bad Good
##           Bad  993 730
##           Good   34 217

```

```

sum(CF20$table*Costs*Prev)

```

```

## [1] -67108.8

```

```

CF30$table

```

```

##           Reference
## Prediction Bad Good
##           Bad  934 531
##           Good   93 416

```

```

sum(CF30$table*Costs*Prev)

```

```

## [1] -28977.6

```

```

CF40$table

```

```

##           Reference
## Prediction Bad Good
##           Bad  859 378
##           Good 168 569

```

```

sum(CF40$table*Costs*Prev)

```

```

## [1] -4497.6

```

```

CF50$table

```

```
##           Reference
## Prediction Bad Good
##           Bad  766  288
##           Good 261  659
```

```
sum(CF50$table*Costs*Prev)
```

```
## [1] 1924.8
```

```
CF60$table
```

```
##           Reference
## Prediction Bad Good
##           Bad  652  180
##           Good 375  767
```

```
sum(CF60$table*Costs*Prev)
```

```
## [1] 9240
```

```
CF70$table
```

```
##           Reference
## Prediction Bad Good
##           Bad  510  103
##           Good 517  844
```

```
sum(CF70$table*Costs*Prev)
```

```
## [1] 4545.6
```

```
CF80$table
```

```
##           Reference
## Prediction Bad Good
##           Bad  318   36
##           Good 709  911
```

```
sum(CF80$table*Costs*Prev)
```

```
## [1] -10708.8
```

```
CF90$table
```

```
##           Reference
## Prediction Bad Good
##           Bad   85    6
##           Good 942  941
```

```
sum(CF90$table*Costs*Prev)
```


```
## [1] -41534.4
```

```
PlotProf<-data.frame(percent_bad_thresh = c(10,20,30,40,50,60,70,80,90),
                      profit = c(sum(CF10$table*Costs*Prev),
                                sum(CF20$table*Costs*Prev),
                                sum(CF30$table*Costs*Prev),
                                sum(CF40$table*Costs*Prev),
                                sum(CF50$table*Costs*Prev),
                                sum(CF60$table*Costs*Prev),
                                sum(CF70$table*Costs*Prev),
                                sum(CF80$table*Costs*Prev),
```



```
sum(CF90$table*Costs*Prev)))

ggplot(PlotProf, aes(y=profit, x=percent_bad_thresh)) +
  geom_point()
```



## Penalized Logistic Regression

```
set.seed(100)
glmGrid <- expand.grid(alpha=c(0, 0.1, 0.2, 0.4, 0.6, 0.8, 1),
                      lambda=seq(0.01, 0.2, length=5))

logreg_penalized_model <- caret::train(x=heloc_train_x,
                                       y=heloc_train_y,
                                       method="glmnet",
                                       metric="ROC",
                                       tuneGrid=glmGrid,
                                       trControl=control)

# bestIndex(logreg_penalized_model)
logreg_penalized_model$results[3:7,1:5]
```

```
##   alpha lambda      ROC      Sens      Spec
## 3   0.0 0.1050 0.7987157 0.7583402 0.7043291
## 4   0.0 0.1525 0.7981792 0.7607732 0.7024836
## 5   0.0 0.2000 0.7976588 0.7619904 0.7011636
## 6   0.1 0.0100 0.8001371 0.7551771 0.7114524
```

```
## 7    0.1 0.0575 0.7992612 0.7551777 0.7043277
```

```
logreg_penalized_modelRoc <- roc_build(logreg_penalized_model)
```

```
## Setting direction: controls < cases
```

Based on the best ROC, has the best specificity, our main metric. This enables us to insure that we only accept best qualified candidates, thus reducing the risk of a loanee defaulting on a \$100,000 loan.

```
# Utilizing Best Model
```

```
testResults$logreg_penalized_model <- stats::predict(logreg_penalized_model,  
                                                    heloc_test_x)
```

```
# confusion matrix
```

```
confusionMatrix(testResults$logreg_penalized_model)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction Bad Good
```

```
##           Bad  773  282
```

```
##           Good 254  665
```

```
##
```

```
##           Accuracy : 0.7285
```

```
##           95% CI : (0.7083, 0.748)
```

```
##           No Information Rate : 0.5203
```

```
##           P-Value [Acc > NIR] : <2e-16
```

```
##
```

```
##           Kappa : 0.4554
```

```
##
```

```
##           McNemar's Test P-Value : 0.2435
```

```
##
```

```
##           Sensitivity : 0.7022
```

```
##           Specificity : 0.7527
```

```
##           Pos Pred Value : 0.7236
```

```
##           Neg Pred Value : 0.7327
```

```
##           Prevalence : 0.4797
```

```
##           Detection Rate : 0.3369
```

```
##           Detection Prevalence : 0.4656
```

```
##           Balanced Accuracy : 0.7274
```

```
##
```

```
##           'Positive' Class : Good
```

```
##
```

## Nonlinear Classification Models

### Flexible Discriminant Analysis

```
# set.seed(100)
```

```
# fdaGrid <- expand.grid(degree=c(1,2),
```

```
#                       nprune=seq(14, 20, 1))
```

```
#
```

```
# set.seed(100)
```

```
# fdaModel <- caret::train(x = heloc_train_x,
```

```
#                       y = heloc_train_y,
```

```
#                       method = "fda",
```

```
#                       metric = "ROC",
```

```
#                               trControl=control,
#                               tuneGrid=fdaGrid)
# # bestIndex(fdaModel)
# fdaModel$results[1:18,1:5]
```

As we see, the best is that of nprune 16 with a specificity of 72.15% (index = 15).

```
fdaGrid <- expand.grid(degree=1,
                      nprune=16)

set.seed(100)
fdaModel <- caret::train(x = heloc_train_x,
                        y = heloc_train_y,
                        method = "fda",
                        metric = "ROC",
                        trControl=control,
                        tuneGrid=fdaGrid)
```

```
## Loading required package: earth
## Loading required package: Formula
## Loading required package: plotmo
## Loading required package: plotrix
## Loading required package: TeachingDemos
fda_modelRoc <- roc_build(fdaModel)
```

```
## Setting direction: controls < cases
testResults$fda_model <- predict(fdaModel,
                                heloc_test_x)

# confusion matrix
confusionMatrix(testResults$fda_model)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction Bad Good
##          Bad  762  265
##          Good 265  682
##
##              Accuracy : 0.7315
##              95% CI : (0.7114, 0.751)
##          No Information Rate : 0.5203
##          P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 0.4621
##
##  Mcnemar's Test P-Value : 1
##
##              Sensitivity : 0.7202
##              Specificity : 0.7420
##          Pos Pred Value : 0.7202
##          Neg Pred Value : 0.7420
##              Prevalence : 0.4797
```

```
##          Detection Rate : 0.3455
##    Detection Prevalence : 0.4797
##          Balanced Accuracy : 0.7311
##
##          'Positive' Class : Good
##
```

## Neural Network

```
# set.seed(100)
# nnetGrid <- expand.grid(size = 1:2,
#                          decay = c(0, 0.1, 0.25, 0.5, 0.75, 1))
#
# nnetModel <- caret::train(x = heloc_train_x,
#                           y = heloc_train_y,
#                           method = "nnet",
#                           tuneGrid = nnetGrid,
#                           metric = "ROC",
#                           trace = FALSE,
#                           maxit = 2000,
#                           trControl = control)
# nnetModel$bestTune$results[1:6,1:5]
```

It appears that the model starts to over fit once the decay goes to 0.1. The nnet model chose size=2 with decay of 2, with nearly identical ROC, Sensitivity, and Specificity scores and thus size 1 with decay 0 is our choice.

```
set.seed(100)
nnetGrid <- expand.grid(size = 1,
                        decay = 0)

nnetModel <- caret::train(x = heloc_train_x,
                          y = heloc_train_y,
                          method = "nnet",
                          tuneGrid = nnetGrid,
                          metric = "ROC",
                          trace = FALSE,
                          maxit = 2000,
                          trControl = control)

nnet_modelRoc <- roc_build(nnetModel)

## Setting direction: controls < cases
testResults$nnet_model <- predict(nnetModel,
                                  heloc_test_x)

# confusion matrix
confusionMatrix(testResults$nnet_model)

## Confusion Matrix and Statistics
##
##          Reference
## Prediction Bad Good
##          Bad  769  291
##          Good 258  656
##
```

```
##           Accuracy : 0.7219
##           95% CI : (0.7015, 0.7416)
##      No Information Rate : 0.5203
##      P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.4421
##
##  McNemar's Test P-Value : 0.172
##
##           Sensitivity : 0.6927
##           Specificity : 0.7488
##      Pos Pred Value : 0.7177
##      Neg Pred Value : 0.7255
##           Prevalence : 0.4797
##      Detection Rate : 0.3323
##      Detection Prevalence : 0.4630
##      Balanced Accuracy : 0.7207
##
##      'Positive' Class : Good
##
```

## Classification Trees

### Boosted Tree

```
# gbmGrid <- expand.grid(interaction.depth = c(2,3),
#                         n.trees = c(1000,2000,3000,4000), #default val = 1000
#                         shrinkage = c(0.01, 0.1),
#                         n.minobsinnode = c(5,10)) # default val = 10
# set.seed(100)
# gbmModel <- caret::train(x = heloc_train_x,
#                           y = heloc_train_y,
#                           method = "gbm",
#                           tuneGrid = gbmGrid,
#                           verbose = FALSE,
#                           metric = "ROC",
#                           trControl= control)
# gbmModel$results
```

```
gbmGrid <- expand.grid(interaction.depth = 2,
                      n.trees = 1000,
                      shrinkage = 0.01,
                      n.minobsinnode = 5)
set.seed(100)
gbmModel <- caret::train(x = heloc_train_x,
                        y = heloc_train_y,
                        method = "gbm",
                        tuneGrid = gbmGrid,
                        verbose = FALSE,
                        metric = "ROC",
                        trControl= control)

gbm_modelRoc <- roc_build(gbmModel)
```

```
## Setting direction: controls < cases
```

```
testResults$gbm_model <- predict(gbmModel,
                                heloc_test_x)

# confusion matrix
confusionMatrix(testResults$gbm_model)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction Bad Good
##      Bad  775  285
##      Good 252  662
##
##           Accuracy : 0.728
##           95% CI : (0.7078, 0.7475)
##      No Information Rate : 0.5203
##      P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.4543
##
##  Mcnemar's Test P-Value : 0.1673
##
##           Sensitivity : 0.6990
##           Specificity : 0.7546
##      Pos Pred Value : 0.7243
##      Neg Pred Value : 0.7311
##           Prevalence : 0.4797
##      Detection Rate : 0.3354
##      Detection Prevalence : 0.4630
##      Balanced Accuracy : 0.7268
##
##      'Positive' Class : Good
##
```

```
# gbmGrid <- expand.grid(interaction.depth = c(2,3),
#                       n.trees = c(1000,2000),
#                       shrinkage = c(0.01, 0.1),
#                       n.minobsinnode = c(5,10))
# set.seed(100)
# gbmMono_model <- caret::train(x = heloc_train_x,
#                               y = heloc_train_y,
#                               method = "gbm",
#                               var.monotone = c(-1,-1,-1,-1,-1,-1,-1,-1,
#                                                1,1,1,1,0,0,-1,0,-1,1,
#                                                0,-1,1),
#                               tuneGrid = gbmGrid,
#                               verbose = FALSE,
#                               metric = "ROC",
#                               trControl= control)
# gbmMono_model$results
```

```
gbmGrid <- expand.grid(interaction.depth = 2,
                     n.trees = 1000,
```

```

shrinkage = 0.01,
n.minobsinnode = 5)
set.seed(100)
gbmMono_model <- caret::train(x = heloc_train_x,
                              y = heloc_train_y,
                              var.monotone = c(-1,-1,-1,-1,-1,-1,-1,-1,
                                                1,1,1,1,0,0,-1,0,-1,1,
                                                0,-1,1),
                              method = "gbm",
                              tuneGrid = gbmGrid,
                              verbose = FALSE,
                              metric = "ROC",
                              trControl= control)

gbmMono_modelRoc <- roc_build(gbmMono_model)

```

### GBM with monotonic constraints

```

## Setting direction: controls < cases
testResults$gbmMono_model <- predict(gbmMono_model,
                                     heloc_test_x)
confusionMatrix(testResults$gbmMono_model)

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction Bad Good
##      Bad  772  281
##      Good 255  666
##
##           Accuracy : 0.7285
##           95% CI : (0.7083, 0.748)
##      No Information Rate : 0.5203
##      P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.4555
##
##  Mcnemar's Test P-Value : 0.2802
##
##           Sensitivity : 0.7033
##           Specificity : 0.7517
##           Pos Pred Value : 0.7231
##           Neg Pred Value : 0.7331
##           Prevalence : 0.4797
##           Detection Rate : 0.3374
##      Detection Prevalence : 0.4666
##           Balanced Accuracy : 0.7275
##
##           'Positive' Class : Good
##

```

## CART

```
# set.seed(100)
# rpart_grid <- expand.grid(cp=c(0.0005, 0.001250, 0.0015, 0.00175, 0.002))
#
# rpart_model <- caret::train(x=heloc_train_x,
#                             y=heloc_train_y,
#                             method="rpart",
#                             metric="ROC",
#                             trControl=control,
#                             tuneGrid = rpart_grid)
#
# testResults$rpart_model <- predict(rpart_model, heloc_test_x)
#
# rpart_model
```

As we see, specificity for this model is one of the worst out of all the models we have.

```
set.seed(100)
rpart_grid <- expand.grid(cp=0.00175)

rpart_model <- caret::train(x=heloc_train_x,
                           y=heloc_train_y,
                           method="rpart",
                           metric="ROC",
                           trControl=control,
                           tuneGrid = rpart_grid)

rpart_modelRoc <- roc_build(rpart_model)

## Setting direction: controls < cases
rpart_model

## CART
##
## 7897 samples
## 21 predictor
## 2 classes: 'Bad', 'Good'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 7107, 7107, 7107, 7108, 7107, 7108, ...
## Resampling results:
##
##   ROC      Sens      Spec
## 0.7558931 0.7551742 0.671846
##
## Tuning parameter 'cp' was held constant at a value of 0.00175

testResults$rpart_model <- predict(rpart_model, heloc_test_x)

confusionMatrix(testResults$rpart_model)

## Confusion Matrix and Statistics
##
```



```
##           Reference
## Prediction Bad Good
##           Bad  748  282
##           Good 279  665
##
##           Accuracy : 0.7158
##           95% CI : (0.6953, 0.7356)
##           No Information Rate : 0.5203
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.4306
##
## Mcnemar's Test P-Value : 0.9327
##
##           Sensitivity : 0.7022
##           Specificity : 0.7283
##           Pos Pred Value : 0.7044
##           Neg Pred Value : 0.7262
##           Prevalence : 0.4797
##           Detection Rate : 0.3369
##           Detection Prevalence : 0.4782
##           Balanced Accuracy : 0.7153
##
##           'Positive' Class : Good
##
```

## Random Forest

```
# set.seed(100)
# rf_grid <- expand.grid(mtry=c(5,10,15))
#
# randomForest_model <- caret::train(x=heloc_train_x,
#                                     y=heloc_train_y,
#                                     method="rf",
#                                     metric="ROC",
#                                     trControl=control,
#                                     tuneGrid = rf_grid)
#
# randomForest_model
```

```
set.seed(100)
rf_grid <- expand.grid(mtry=5)

randomForest_model <- caret::train(x=heloc_train_x,
                                   y=heloc_train_y,
                                   method="rf",
                                   metric="ROC",
                                   trControl=control,
                                   tuneGrid = rf_grid)

randomForest_modelRoc <- roc_build(randomForest_model)
```

```
## Setting direction: controls < cases
```

```

randomForest_model

## Random Forest
##
## 7897 samples
## 21 predictor
## 2 classes: 'Bad', 'Good'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 7107, 7107, 7107, 7108, 7107, 7108, ...
## Resampling results:
##
## ROC      Sens      Spec
## 0.7951729 0.7736746 0.6847978
##
## Tuning parameter 'mtry' was held constant at a value of 5
testResults$randomForest_model <- predict(randomForest_model, heloc_test_x)

confusionMatrix(testResults$randomForest_model)

## Confusion Matrix and Statistics
##
##           Reference
## Prediction Bad Good
##      Bad 790 291
##      Good 237 656
##
##           Accuracy : 0.7325
##           95% CI : (0.7124, 0.7519)
##      No Information Rate : 0.5203
##      P-Value [Acc > NIR] : < 2e-16
##
##           Kappa : 0.463
##
##  Mcnemar's Test P-Value : 0.02108
##
##           Sensitivity : 0.6927
##           Specificity : 0.7692
##           Pos Pred Value : 0.7346
##           Neg Pred Value : 0.7308
##           Prevalence : 0.4797
##           Detection Rate : 0.3323
##      Detection Prevalence : 0.4524
##           Balanced Accuracy : 0.7310
##
##           'Positive' Class : Good
##

```

## Result Discussion

```

plot(lda_modelRoc, type='s', col='antiquewhite4', legacy.axes=TRUE)
plot(logreg_modelRoc, type='s', col='aquamarine3', legacy.axes=TRUE, add=TRUE)

```


```

plot(logreg_penalized_modelRoc, type='s', col='blue', legacy.axes=TRUE, add=TRUE)
plot(fda_modelRoc, type='s', col='blueviolet', legacy.axes=TRUE, add=TRUE)
plot(nnet_modelRoc, type='s', col='brown', legacy.axes=TRUE, add=TRUE)
plot(gbm_modelRoc, type='s', col='cadetblue', legacy.axes=TRUE, add=TRUE)
plot(gbmMono_modelRoc, type='s', col='red', legacy.axes=TRUE, add=TRUE)
plot(rpart_modelRoc, type='s', col='chartreuse', legacy.axes=TRUE, add=TRUE)
plot(randomForest_modelRoc, type='s', col='cornflowerblue', legacy.axes=TRUE, add=TRUE)

legend_ <- c('LDA', 'LR', 'Penalized LR', 'FDA', 'NNET', 'GBM', 'GBM+mono', 'RPART', 'RandomForest')
colors_ <- c('antiquewhite4',
             'aquamarine3',
             'blue',
             'blueviolet',
             'brown',
             'cadetblue',
             'red',
             'chartreuse',
             'cornflowerblue')
legend('bottomright', legend=legend_,
      col=colors_, lwd=2)
title(main = 'Compare ROC curves from different models', outer = TRUE)

```

### Compare ROC curves from different models



```

# gather all model AUCs in 1 list
aucs <- c(lda_modelRoc$auc,
          logreg_modelRoc$auc,
          logreg_penalized_modelRoc$auc,
          fda_modelRoc$auc,

```

```

nnet_modelRoc$auc,
gbm_modelRoc$auc,
gbmMono_modelRoc$auc,
rpart_modelRoc$auc,
randomForest_modelRoc$auc)

scoreboard <- modelScoreBoard(testResults)
# add AUC list as a column
scoreboard$AUC <- aucs
scoreboard

```

##	Accuracy	Precision	Sensitivity	Specificity	F1
## lda_model	0.7228977	0.7178649	0.6958817	0.7478092	0.7067024
## log_reg_model	0.7218845	0.7163043	0.6958817	0.7458617	0.7059454
## logreg_penalized_model	0.7284701	0.7236126	0.7022175	0.7526777	0.7127546
## fda_model	0.7315096	0.7201690	0.7201690	0.7419669	0.7201690
## nnet_model	0.7218845	0.7177243	0.6927138	0.7487829	0.7049973
## gbm_model	0.7279635	0.7242888	0.6990496	0.7546251	0.7114455
## gbmMono_model	0.7284701	0.7231270	0.7032735	0.7517040	0.7130621
## rpart_model	0.7158055	0.7044492	0.7022175	0.7283350	0.7033316
## randomForest_model	0.7325228	0.7346025	0.6927138	0.7692308	0.7130435
##	AUC				
## lda_model	0.7998810				
## log_reg_model	0.8002390				
## logreg_penalized_model	0.7820838				
## fda_model	0.8032575				
## nnet_model	0.8001313				
## gbm_model	0.7998948				
## gbmMono_model	0.8003082				
## rpart_model	0.7554376				
## randomForest_model	0.7950744				

## REFERENCES:

Kuhn, M., & Johnson, K. (2013). Applied Predictive Modeling. New York: Springer.